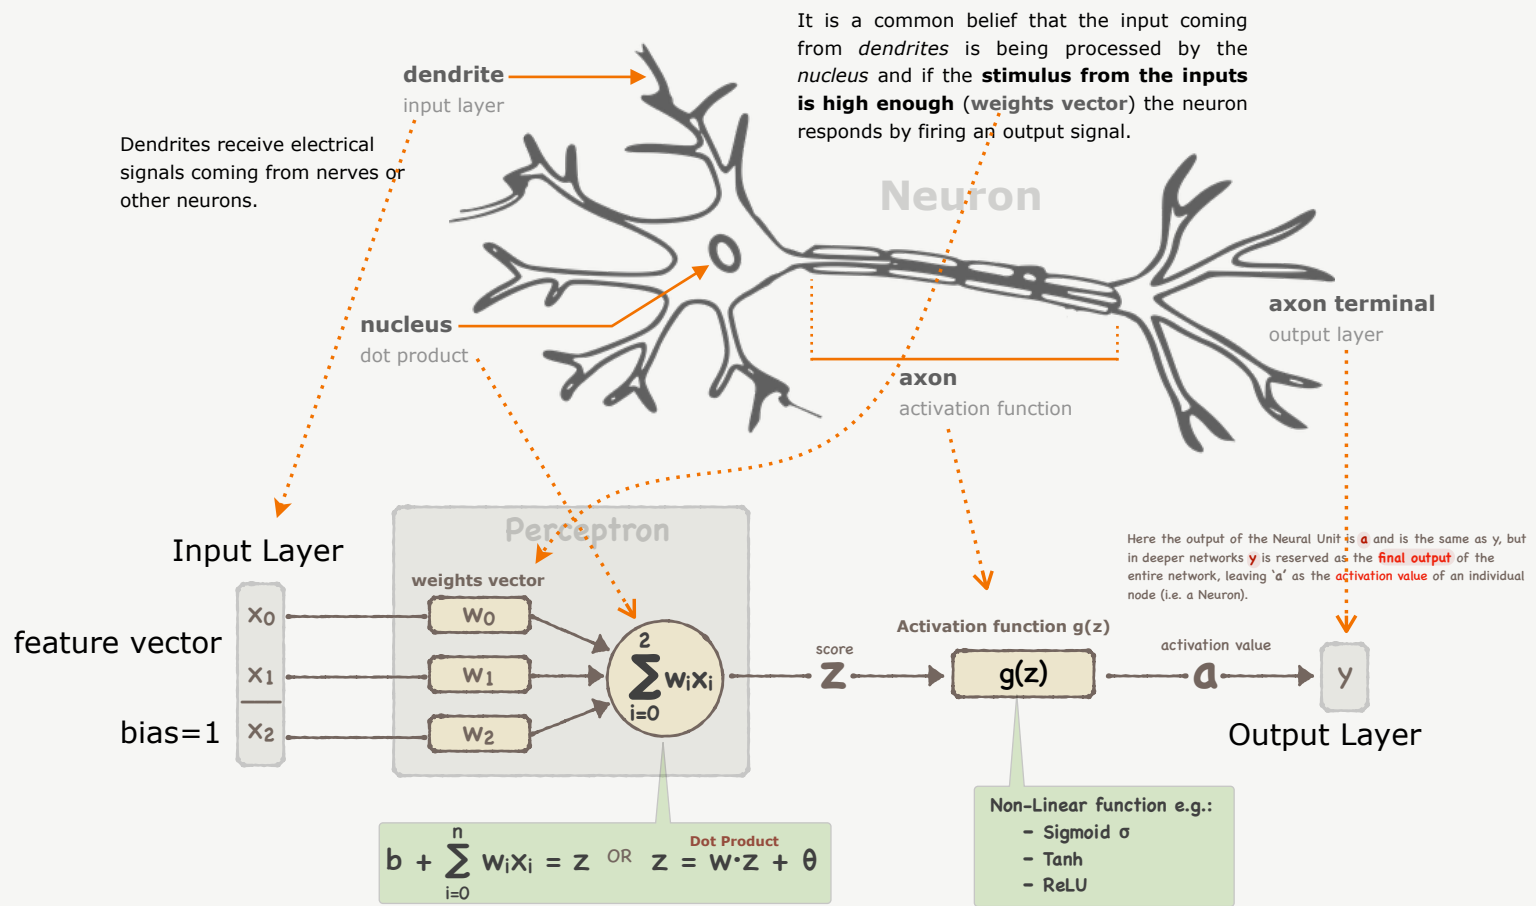


Modeling a Neuron

The human brain is a biological network made out of neurons. Each neuron is said to perform a very moderate cognitive function but collectively contributing to much larger cognitive tasks.


[illegible]

"A journey of a thousand miles begins with a single step" a Chinese proverb.

Building a neural network may seem like a task too large to overcome. When facing a difficult task, it is a sound strategy to break it down into its constituent components and tackle those instead. Let’s use our neuron model to act like a logic gate.

AND Gate


Given four points (0,0), (0,1), (1,0), (1,1) it is quite easy to separate them in a linear fashion i.e. using a straight line.



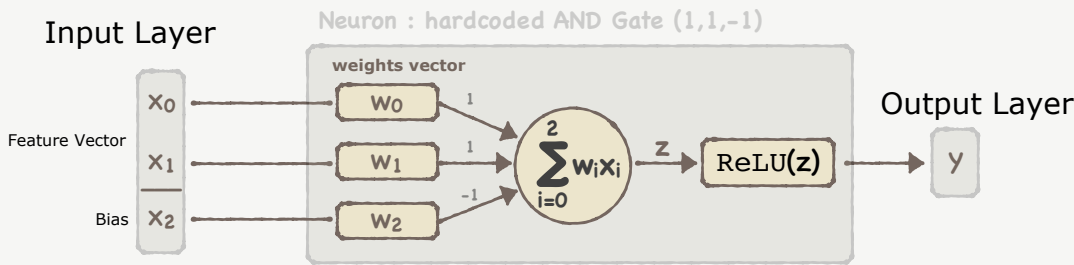
the AND Gate looks like this:

AND Gate		
x ₀	x ₁	output
0	0	0
0	1	0
1	0	0
1	1	1

So, when inputting e.g. (0,0), one must expect that the Perceptron model (given the appropriate set of weights) will output (0).




Let’s model a Perceptron model that does just that. Also let’s use ReLU function as model’s activation function.



Below is the working definition of **ReLU** and its derivative:

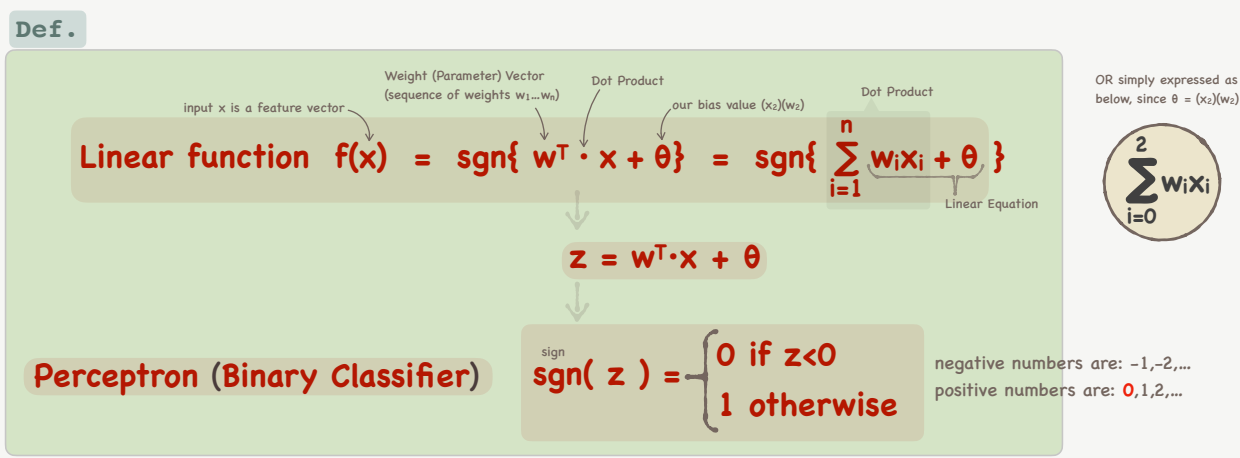
$$\text{ReLU}(z) = \begin{cases} 0, & \text{if } z < 0 \\ z, & \text{otherwise} \end{cases}$$
$$\frac{d}{dz} \text{ReLU}(z) = \begin{cases} 0, & \text{if } z < 0 \\ 1, & \text{otherwise} \end{cases}$$



Note that the derivative of ReLU is the **unit step function** which is NOT a non-linear function—however, in our case, it does NOT have to be since our data is linearly separable. Let’s call our unit step function as ‘**sgn**’ (short for + or – sign).

The slope-intercept equation of a line with slope **m** and y-intercept **b** is: **y = mx + b** where in our situation:

- x** : is the input **x_i**, for 0 ≤ i ≤ 1
- m** : is the weight **w_i** (applied to input **x_i**)
- b** : is the bias term **x₂**
- y** : is the decision boundary (i.e. the line that separates our input data)



First, I tested the model using the hardcoded set of weights that are guaranteed to produce the same output as the AND Gate, i.e. (w₀,w₁,w₂) = (1,1,-1).

given input (x₀,x₁) and bias x₂ = 1
(x₀,x₁,x₂)*(w₀,w₁,w₂) = (x₀w₀)+(x₁w₁)+(x₂w₂) = z → **sgn(z)** = output
(1,1,1)*(1,1,-1) = (1+1+(-1)) = 1 → **sgn(1)** = 1
(1,0,1)*(1,1,-1) = (1+0+(-1)) = 0 → **sgn(0)** = 0
(0,1,1)*(1,1,-1) = (0+1+(-1)) = 0 → **sgn(0)** = 0
(0,0,1)*(1,1,-1) = (0+0+(-1)) = -1 → **sgn(-1)** = 0

Below is the program's output given the hardcoded weights (1,1,-1)

```
AND GATE:
0 AND 0 : 0
0 AND 1 : 0
1 AND 0 : 0
1 AND 1 : 1
Program ended with exit code: 0
```

After running the network using backpropagation (i.e. making the model learn its own set of weights), it just happened so that the model, given ReLU as its activation function, learned the same exact set of weights as I provided in the hardcoded test above. I run the model for 100 epochs (iterations). Notice how rapidly the mean squared error (MSE) (i.e. the difference between the model’s output and the ground truth) is diminishing to ≈ zero.

```
Epoch: 0 MSE: 0.25
Epoch: 10 MSE: 0.00989249
Epoch: 20 MSE: 3.28619e-05
Epoch: 30 MSE: 1.04298e-07
Epoch: 40 MSE: 3.30755e-10
Epoch: 50 MSE: 1.04889e-12
Epoch: 60 MSE: 3.32627e-15
Epoch: 70 MSE: 1.05483e-17
Epoch: 80 MSE: 3.34509e-20
Epoch: 90 MSE: 1.06079e-22

Below are the trained/learned weights:
Layer-1 : Input Layer
Layer-2 : Neuron-1: [ 1 1 -1 ]

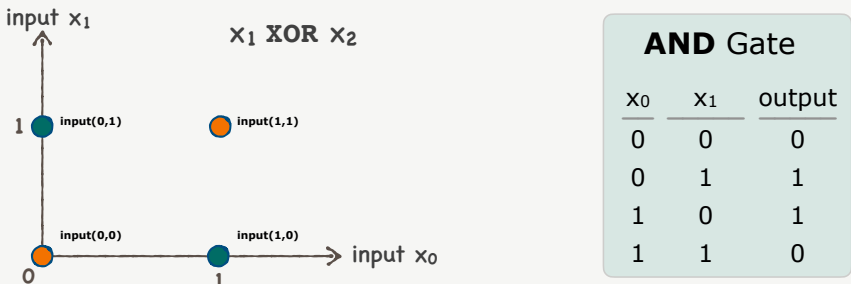
AND GATE:
0 AND 0 : 0
0 AND 1 : 0
1 AND 0 : 0
1 AND 1 : 1
Program ended with exit code: 0
```

The purpose of using non-linear activation functions is to be able to solve non-linearly separable problems. Simulating XOR Gate is such a task.

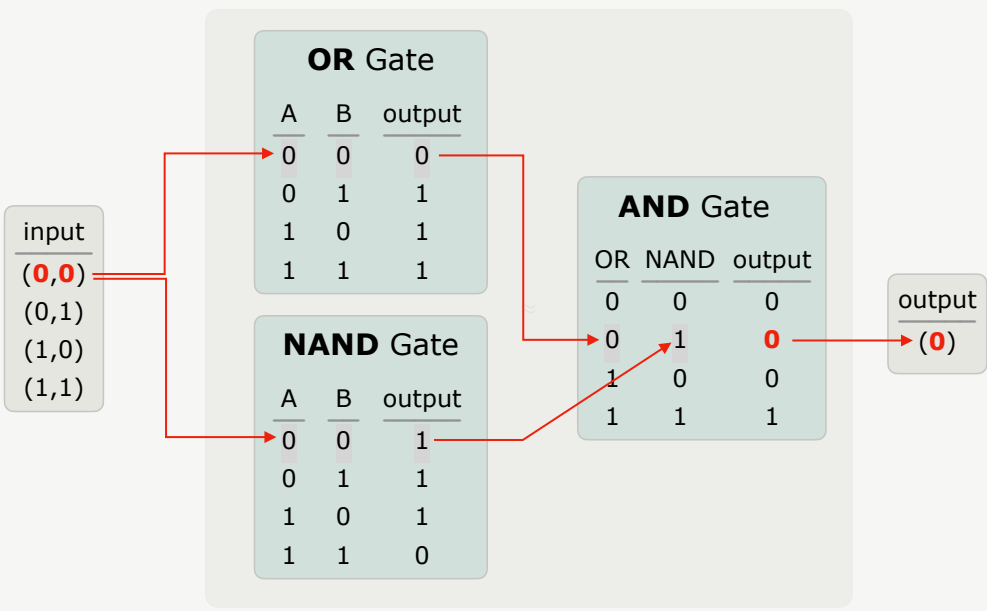
XOR Gate

Given four points (0,0), (0,1), (1,0), (1,1) it is NOT possible to separate them in a linear fashion i.e. using a straight line (go ahead try for yourself).

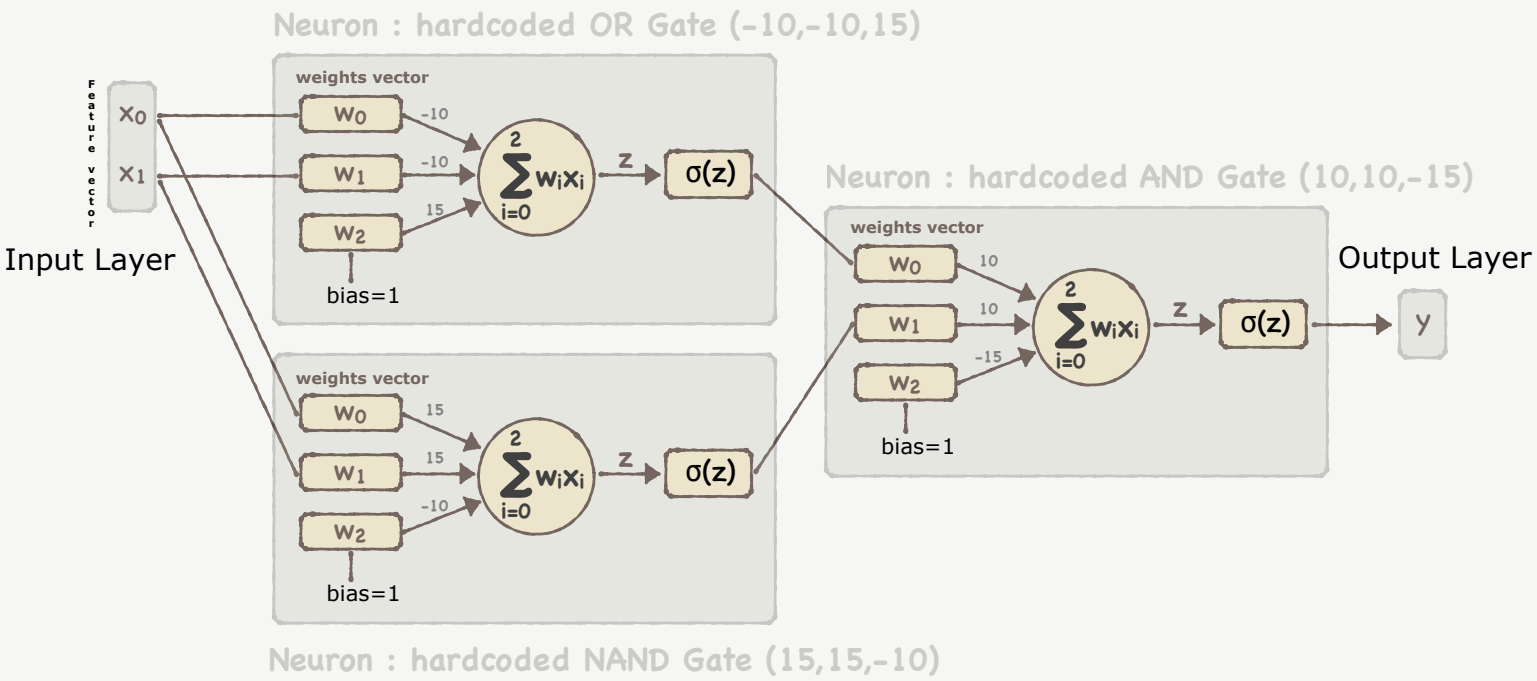
the XOR Gate looks like this:



It is a common approach to implement the XOR gate as a composition of **OR**, **NAND** and **AND** gates. For example given input (0,0) we expect the output (given the appropriate set of weights for each gate) to be (0).



Let’s model a Perceptron model to do just that. But this time, let’s use either Sigmoid or TanH as model’s activation function. YES—we’ll need three neurons.

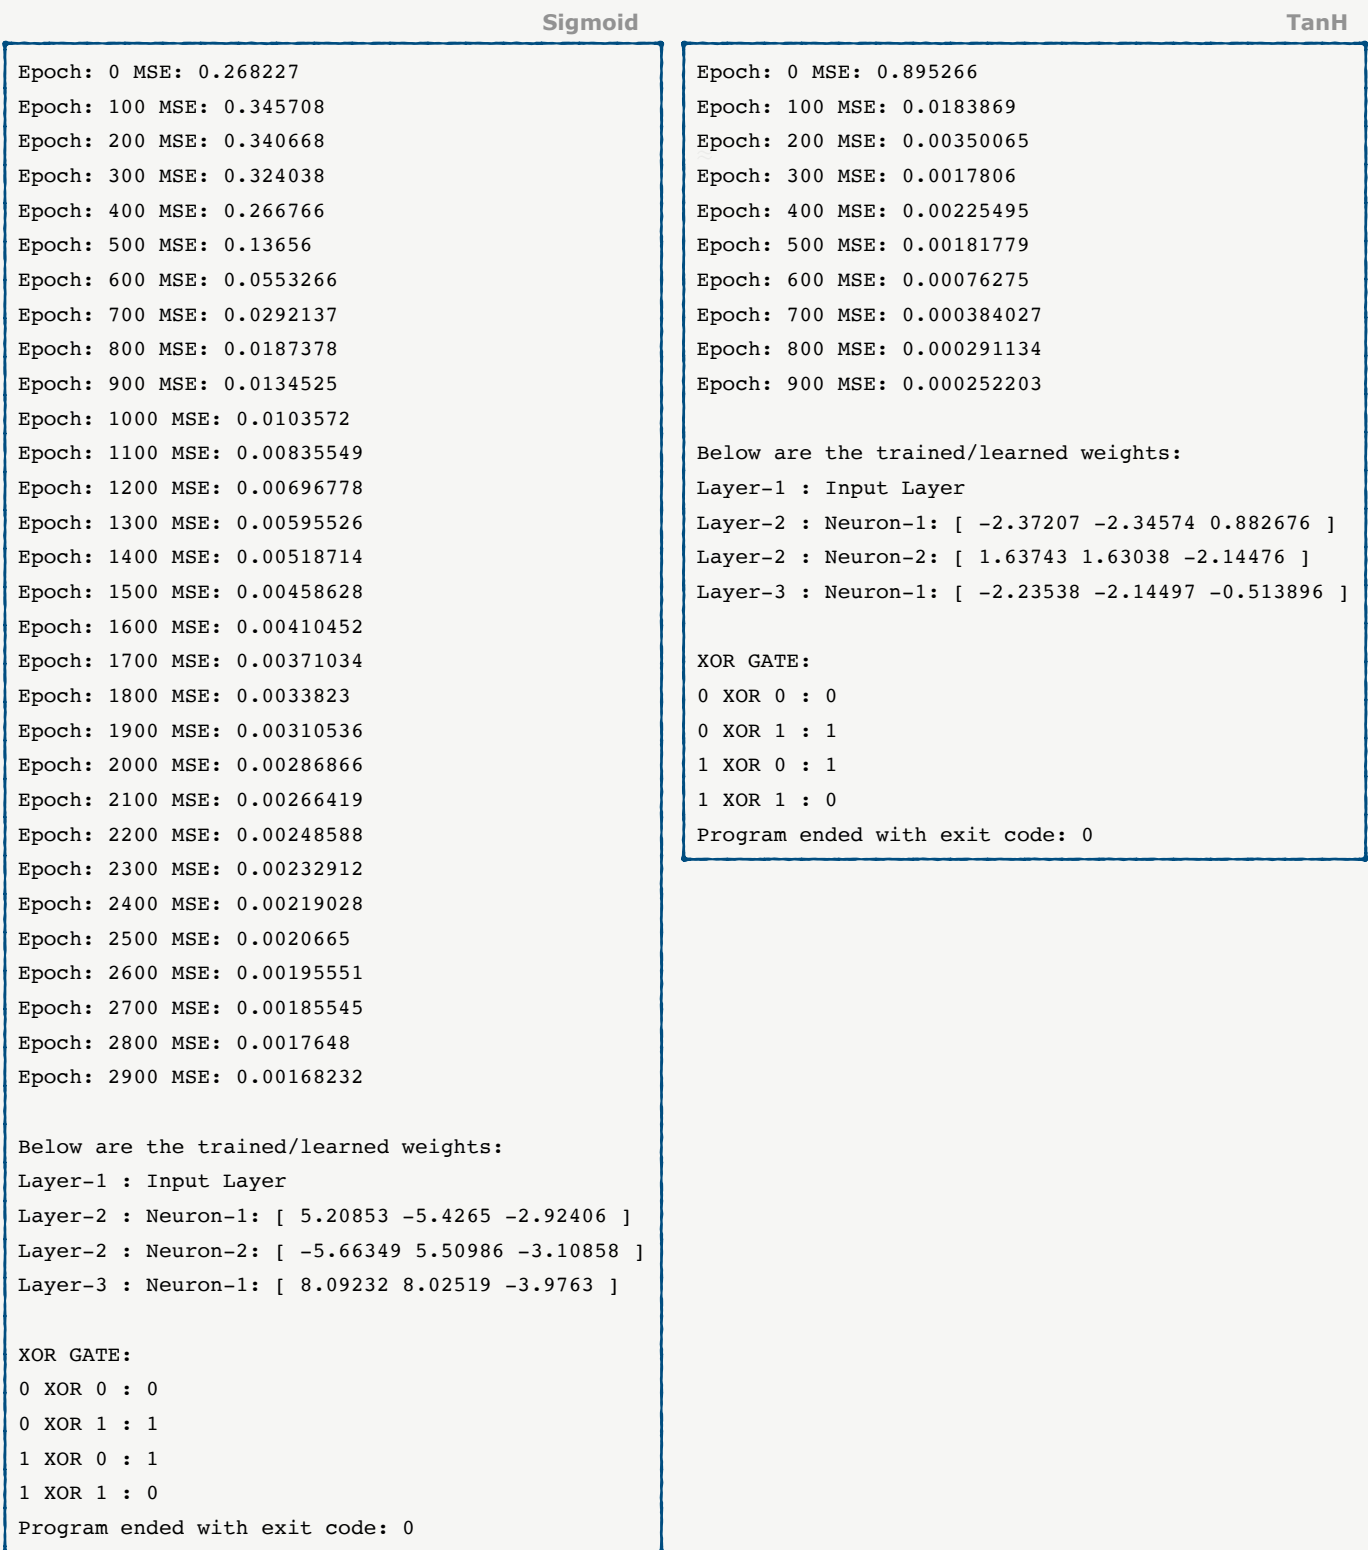


Given the hardcoded set of weights, the model did perform as expected:

```
Layer-1 : Input Layer
Layer-2 : Neuron-1: [ -10 -10 15 ]
Layer-2 : Neuron-2: [ 15 15 -10 ]
Layer-3 : Neuron-1: [ 10 10 -15 ]

XOR GATE:
0 XOR 0 : 0
0 XOR 1 : 1
1 XOR 0 : 1
1 XOR 1 : 0
```

Next I run the model for 3000 epochs using Sigmoid and 1000 epochs using TanH. Below are the result of training/learning the necessary weights; as well as the resulting output of the model while using those weights.



Using the ReLU whose derivative is NOT a non-linear function, the model could not correctly simulate the XOR gate.

